

# Self-Driving Cars

## Exercise 0 - Introduction

Kashyap Chitta

Autonomous Vision Group  
MPI-IS / University of Tübingen

October 18, 2019



University of Tübingen  
MPI for Intelligent Systems  

---

Autonomous Vision Group



# Outline

- ▶ Setup for Exercises
- ▶ PyTorch
- ▶ OpenAI Gym
- ▶ Exercise 0

## Setup for Exercises

# ILIAS

- ▶ We organize the exercises using the ILIAS system  
`https://ovidius.uni-tuebingen.de/ilias3`
- ▶ Exercise sheets will be available in the ILIAS system. Please be aware of the **submission deadline**.
- ▶ You are eligible to finish the homework within a group up to 2 people, but **each person must submit a solution**.
- ▶ If you have any questions, please ask at the **forum** on ILIAS.
- ▶ TA Email IDs: `eshed.ohn-bar@tue.mpg.de`, `kashyap.chitta@tue.mpg.de`,  
`katja.schwarz@tue.mpg.de`.

# TCML cluster

- ▶ You are eligible to use the Training Center for Machine Learning (TCML) cluster for exercises in this lecture.
- ▶ We will create accounts for you on the cluster. Each group shares an account.  
**You are not supposed to apply for the account by yourself.**
- ▶ Cluster has a master node to launch jobs, and 40 compute nodes to execute jobs.
- ▶ Compute nodes are allocated based on a queuing system, so please **start the assignments early** if you want to use this resource.

# TCML cluster

- ▶ Login to master:

```
ssh username@tcml-master01.uni-tuebingen.de
```

- ▶ To access compute nodes: create an `.sbatch` file

- ▶ Please find more information/instructions about the cluster below:

```
https://docs.google.com/document/d/
```

```
1AgtLy28VVZaPe79Tw0b9jjC4F1KVzffb8y1vZoURZE8/edit?usp=sharing.
```

PyTorch

# PyTorch

- ▶ What is PyTorch?

A Python-based scientific computing package for Deep Learning.

- ▶ Why PyTorch?

Beginner friendly, well documented, good for fast development.

- ▶ How to install?

`https://pytorch.org/get-started/locally/`

- ▶ This tutorial is for **PyTorch 1.3**



# Basic Computations

## Tensor

### ► Construct a Tensor

```
x = torch.zeros(8, 3)
print(x)
```

```
tensor([[0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.]])
```

```
x = torch.rand(8, 3)
print(x)
```

```
tensor([[0.0984, 0.3671, 0.2543],
        [0.5016, 0.8017, 0.0918],
        [0.5081, 0.6020, 0.0867],
        [0.5313, 0.4571, 0.1624],
        [0.4231, 0.3993, 0.2713],
        [0.8978, 0.6039, 0.8519],
        [0.4829, 0.6648, 0.8295],
        [0.9060, 0.2132, 0.4110]])
```

# Basic Computations

## Operations

- Multiple syntaxes, e.g. Addition

```
y = torch.rand(8, 3)
```

```
print(x + y)
```

```
print(torch.add(x, y))
```

```
# providing an output tensor as argument
```

```
result = torch.empty(8, 3)
```

```
torch.add(x, y, out=result)
```

```
print(result)
```

```
# adds x to y
```

```
y.add_(x)
```

```
print(y)
```

- Other operations including transposing, indexing, slicing, linear algebra etc. at <https://pytorch.org/docs/stable/torch.html>

# Basic Computations

## Bridge to Numpy

- ▶ PyTorch → Numpy

```
a = torch.ones(5)
b = a.numpy()
```

- ▶ Numpy → PyTorch

```
a = np.ones(5)
b = torch.from_numpy(a)
```

- ▶ Tensors can only be converted to Numpy when they are on CPU

# Basic Computations

## Difference to Numpy

### ► GPU acceleration

```
if torch.cuda.is_available():  
    x = x.cuda()  
    y = y.cuda()  
    z = x + y  
    print(z)  
    print(z.cpu())          # ‘.to’ can change dtype
```

```
tensor([2.9218], device='cuda:0')  
tensor([2.9218], dtype=torch.float64)
```

# Basic Computations

## Difference to Numpy

- ▶ GPU acceleration
- ▶ Automatic differentiation for all operations on Tensors

```
x = torch.ones(2, 2, requires_grad=True)
y = x + 2
z = y * y * 3
out = z.mean()


out.backward()

print(x.grad)
```

```
tensor([[4.5000, 4.5000],
        [4.5000, 4.5000]])
```


# Neural Networks

- ▶ Model → `torch.nn.Module`
- ▶ Loss → `torch.nn.Module.Loss`
- ▶ Optimizer → `torch.optim`
- ▶ Data → `torch.utils.data`



# Neural Networks

## Model



# Neural Networks

## Model

- ▶ Define a model as a class that inherits from `torch.nn.Module`

```
class Net(nn.Module):
```

- ▶ Define layers in the `__init__()` method

```
    def __init__(self):  
        ...
```

- ▶ Define computation flow given an input `x` in the `forward()` method

```
    def forward(self, x):  
        ...
```

- ▶ `backward()` is automatically defined



# Neural Networks

## Model

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        return x
```

# Neural Networks

## Model

- ▶ PyTorch contains a branch of standard layers which are also subclasses of `torch.nn.Module`:

- ▶ Convolution layers `nn.Conv2d( $C_{in}$ ,  $C_{out}$ ,  $K$ )`
- ▶ Pooling layers `nn.MaxPool2d( $K$ )`
- ▶ Non-linear activations `nn.ReLU()`
- ▶ Normalization layers `nn.BatchNorm2d( $N$ )`
- ▶ Linear layers `nn.Linear( $C_{in}$ ,  $C_{out}$ )`
- ▶ ... `...`

- ▶ It is also easy implement custom layers:

<https://pytorch.org/docs/stable/notes/extending.html#extending-torch-nn>

# Neural Networks

## Loss

- ▶ Loss function returns a non-negative value  $J$  measuring the distance between network estimation and the ground truth
- ▶ PyTorch contains a branch of loss functions which are also subclasses of `torch.nn.Module`:
  - ▶ L1Loss
  - ▶ MSELoss
  - ▶ CrossEntropyLoss
  - ▶ NLLLoss
  - ▶ SmoothL1Loss
  - ▶ ...

# Neural Networks

## Loss

- Example of using a loss function

```
loss = nn.CrossEntropyLoss()  
input = torch.randn(3, 5, requires_grad=True)  
target = torch.empty(3, dtype=torch.long).random_(5)  
output = loss(input, target)  
output.backward()
```

# Neural Networks

## Optimizer

- ▶ Optimizer decides how to update the parameters in the model, e.g.

$$\theta = \theta - \eta \nabla J(\theta)$$

- ▶ PyTorch implements a set of optimization algorithms in `torch.optim`:
  - ▶ SGD
  - ▶ SGD + Momentum
  - ▶ Adam
  - ▶ ...

# Neural Networks

## Optimizer

- ▶ 1. Construct an Optimizer

```
optimizer = optim.SGD(model.parameters(), lr = 0.01, momentum=0.9)
```

- ▶ 2. Take an optimization step for every batch/sample

```
for input, target in dataset:
    # clear saved gradients before computing gradient for the new batch
    optimizer.zero_grad()

    output = model(input)
    loss = loss_fn(output, target)

    loss.backward()

    # update parameters in model
    optimizer.step()
```

# Neural Networks

## Data

- ▶ PyTorch provides `Dataset`, `DataLoader` in `torch.utils.data` that allows batching data, shuffling data and load data with multiple processes. Good tutorial at [https://pytorch.org/tutorials/beginner/data\\_loading\\_tutorial.html](https://pytorch.org/tutorials/beginner/data_loading_tutorial.html)
- ▶ For small scale of dataset it is fine to implement your own data loader.

# Neural Networks

## Saving Models

- Save/Load `state_dict` (Recommended)

```
# save
torch.save(model.state_dict(), PATH)

# load
model = TheModelClass(*args, **kwargs)
model.load_state_dict(torch.load(PATH))
model.eval()
```

- Save/Load entire model

```
# save
torch.save(model, PATH)

# load
# Model class must be defined somewhere
model = torch.load(PATH)
model.eval()
```



# References

- ▶ PyTorch official tutorials: <https://pytorch.org/tutorials/>
- ▶ Stanford Course on Deep Learning for Computer Vision:  
[http://cs231n.stanford.edu/slides/2018/cs231n\\_2018\\_lecture08.pdf](http://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture08.pdf)
- ▶ NTU Machine Learning Course: <https://www.slideshare.net/lymanblueLin/pytorch-tutorial-for-ntu-machine-learning-course-2017>
- ▶ PyTorch tutorial with code examples:  
<https://github.com/MorvanZhou/PyTorch-Tutorial>

OpenAI Gym

# OpenAI Gym

- ▶ What is OpenAI Gym?

A python based toolkit for developing and comparing RL algorithms.


- ▶ Why OpenAI Gym?

Standardization of environments/benchmarks for RL algorithms.


- ▶ How to install?

We provide a custom package with instructions.


# Examples



(a) Atari




(b) Box2D




(c) Control



(d) MuJoCo




(e) Robotics



(f) Grid Worlds


# Reinforcement Learning

- ▶ 1. Create an Environment and initialize the State
- ▶ 2. Iteratively take an Action and observe the State, with the goal to maximize the Reward



# Reinforcement Learning

- ▶ 1. Create an Environment and initialize the State
- ▶ 2. Iteratively take an Action and observe the State, with the goal to maximize the Reward



# Work flow

```
import gym
env = gym.make('CarRacing-v0')
env.reset()
for _ in range(1000):

    env.render()

    # take a random action
    env.step(env.action_space.sample())
```

# Work flow

## Create an Environment

- ▶ Each gym environment has a unique name
- ▶ To create an environment from the name use the

```
env = gym.make(env_name)
```

- ▶ For example, to create a CarRacing environment:

```
env = gym.make('CarRacing-v0')
```



# Work flow

## **Initialize State**

- ▶ Used to reinitialize a new episode
- ▶ Returns the initial state

```
init_state = env.reset()
```

# Work flow

## Take an action

- ▶ Performs the specified action and returns the resulting state
- ▶ The main method your agent interacts with

```
step(action) -> (next_state ,  
                 reward ,  
                 is_terminal ,  
                 debug_info)
```

# Work flow


## **Take an observation: Render**

- ▶ Optional method
- ▶ Used to display the state of your environment
- ▶ Useful for debugging and qualitatively comparing different agent policies

```
env.render()
```

# Car Racing

- Randomly generated tracks



# Car Racing

- ▶ `action_space`: three continuous values, including steer, gas, brake


```
self.action_space = spaces.Box( low=np.array([-1,0,0]),  
                                high=np.array([+1,+1,+1]),  
                                dtype=np.float32)
```

- ▶ `observation_space`: color image

```
self.observation_space = spaces.Box(low=0, high=255,  
                                     shape=(STATE_H, STATE_W, 3),  
                                     dtype=np.uint8)
```

- ▶ reward:  $R = N_{visited\_tile} * \frac{1000}{N_{all\_tile}} - N_{frame} * 0.1$

# Car Racing



# Car Racing



► Reward

# Car Racing



- ▶ Reward
- ▶ Car speed



# Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed

# Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed
- ▶ Joint angle

# Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed
- ▶ Joint angle
- ▶ Angular Velocity

# Car Racing

- ▶ We will create a leaderboard for the exercises
- ▶ Evaluation metrics:
  - ▶  $R = N_{visited\_tile} * \frac{1000}{N_{all\_tile}}$  given a fixed  $N_{frame}$
  - ▶  $R = R - 100$  if the car get too far away from the track
- ▶ Submit your code and we will evaluate it on our server

# References

- ▶ Official document: <https://gym.openai.com/docs>
- ▶ Source code: <https://github.com/openai/gym>
- ▶ [https://katefvision.github.io/10703\\_openai\\_gym\\_recitation.pdf](https://katefvision.github.io/10703_openai_gym_recitation.pdf)

## Exercise 0

# Exercise 0

## Install PyTorch locally

- Recommended to install with Anaconda

PyTorch Build	Stable		Preview		
Your OS	Linux		Mac		Windows
Package	Conda		Pip	LibTorch	Source
Language	Python 2.7	Python 3.5	Python 3.6	Python 3.7	C++
CUDA	8.0	9.0	9.2	None	
Run this Command:	conda install pytorch torchvision -c pytorch				

# Exercise 0

## Install OpenAI Gym locally

- ▶ Python 3.5+
- ▶ Download `sdc_gym.zip` from **ILIAS**, unzip and enter the folder
- ▶ Install the box2d package

```
pip3 install -e '[box2d]'
```

- ▶ Please use the code we provided as there are some modifications compared to the official version.



# Exercise 0

## Work on Cluster

- ▶ Download the Singularity image and copy it to your home directories on the cluster: `https://owncloud.tuebingen.mpg.de/index.php/s/CbGdQrCfcP4EFA/download`  
It is an environment that contains everything you need for the exercises of this lecture such as PyTorch, OpenAI Gym.
- ▶ Run the code under the Singularity environment:

```
singularity exec ~/sdc_gym.simg python your_python_file.py
```

Questions?