

Self-Driving Cars

Exercise 1 - Imitation Learning

Eshed Ohn-Bar

Autonomous Vision Group
MPI-IS / University of Tübingen

October 25, 2019



University of Tübingen
MPI for Intelligent Systems

Autonomous Vision Group



Exercise Setup

Download

- ▶ `01_imitation_learning_exercise.pdf`
- ▶ `main.py`
- ▶ `network.py`, `training.py`, `imitations.py`
- ▶ `data`, a folder with imitations

Exercise Setup

Download

- ▶ `01_imitation_learning_exercise.pdf`
- ▶ `main.py`
- ▶ `network.py`, `training.py`, `imitations.py`
- ▶ `data`, a folder with imitations

Submit

- ▶ a report as a `.pdf` file, up to 3 pages
- ▶ your code: `network.py`, `training.py`, `imitations.py` as a `.zip` file
- ▶ your pre-trained model as a `.t7` file

Deadline: **Wed, 21. November 2018 - 21:00**

Exercise Setup

Do's

- ▶ comment your code
- ▶ use docstrings
- ▶ use self-explanatory variable names
- ▶ structure your code well

Exercise Setup

Do's

- ▶ comment your code
- ▶ use docstrings
- ▶ use self-explanatory variable names
- ▶ structure your code well

Do not's

- ▶ change `main.py`, especially `calculate_score_for_leaderboard()`
- ▶ install more packages
- ▶ change the gym environment

Imitation Learning

Behavioral Cloning

Imitation Learning

Components:

- ▶ State: $s \in \mathcal{S}$ may be partially observed (e.g., game screen)
- ▶ Action: $a \in \mathcal{A}$ may be discrete or continuous (e.g., turn angle, speed)
- ▶ Policy: $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ we want to learn the policy parameters θ
- ▶ Optimal action: $a^* \in \mathcal{A}$ provided by expert demonstrator
- ▶ Optimal policy: $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ provided by expert demonstrator
- ▶ State dynamics: $P(s_{i+1}|s_i, a_i)$ simulator, typically not known to policy
Often deterministic: $s_{i+1} = T(s_i, a_i)$ deterministic mapping
- ▶ Rollout: Given s_0 , sequentially execute $a_i = \pi_\theta(s_i)$ & sample $s_{i+1} \sim P(s_{i+1}|s_i, a_i)$
yields trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$
- ▶ Loss function: $\mathcal{L}(a^*, a)$ loss of action a given optimal action a^*

Imitation Learning

Components:

- ▶ State: $s \in \mathcal{S}$ may be partially observed (e.g., game screen)
- ▶ Action: $a \in \mathcal{A}$ may be discrete or continuous (e.g., turn angle, speed)
- ▶ Policy: $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ we want to learn the policy parameters θ
- ▶ Optimal action: $a^* \in \mathcal{A}$ provided by expert demonstrator
- ▶ Optimal policy: $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ provided by expert demonstrator
- ▶ State dynamics: $P(s_{i+1}|s_i, a_i)$ simulator, typically not known to policy
Often deterministic: $s_{i+1} = T(s_i, a_i)$ deterministic mapping
- ▶ Rollout: Given s_0 , sequentially execute $a_i = \pi_\theta(s_i)$ & sample $s_{i+1} \sim P(s_{i+1}|s_i, a_i)$
yields trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$
- ▶ Loss function: $\mathcal{L}(a^*, a)$ loss of action a given optimal action a^*

1.1

Network Design

a) Load imitations

```
7  def load_imitations(data_folder):
8      """
9      1.1 a)
10     Given the folder containing the expert imitations, the data gets loaded and
11     stored it in two lists: observations and actions.
12     |
13     |      N = number of (observation, action) - pairs
14     |      data_folder: python string, the path to the folder containing the
15     |                     observation_%05d.npy and action_%05d.npy files
16     |
17     |      return:
18     |      observations: python list of N numpy.ndarrays of size (96, 96, 3)
19     |      actions:      python list of N numpy.ndarrays of size 3
20     |      """
21     pass
```


b) Understand training

```
8  def train(data_folder, trained_network_file):
9      """
10     Function for training the network.
11     """
12     infer_action = ClassificationNetwork()
13     optimizer = torch.optim.Adam(infer_action.parameters(), lr=1e-2)
14     observations, actions = load_imitations(data_folder)
15     observations = [torch.Tensor(observation) for observation in observations]
16     actions = [torch.Tensor(action) for action in actions]
17
18     batches = [batch for batch in zip(observations,
19                                     infer_action.actions_to_classes(actions))]
20     gpu = torch.device('cuda')
21
22     nr_epochs = 100
23     batch_size = 64
24     nr_of_classes = 0 # needs to be changed
25     start_time = time.time()
26
27     for epoch in range(nr_epochs):
28         random.shuffle(batches)
```

c) Classification Network

- ▶ expert imitations
action = [steer, gas, brake]
e.g. [1., 0., 0.8]
- ▶ define **action-classes**
 - ▶ {steer_left}
 - ▶ {}
 - ▶ {steer_right, brake}
 - ▶ {gas}


⇒ map [1., 0., 0.8] → ?



c) Classification Network


- ▶ expert imitations
action = [steer, gas, brake]
e.g. [1., 0., 0.8]
- ▶ define **action-classes**
 - ▶ {steer_left}
 - ▶ {}
 - ▶ {steer_right, brake}
 - ▶ {gas}

⇒ map [1., 0., 0.8] → [0, 0, 1, 0]



c) Classification Network

- ▶ `actions_to_classes`
expert imitations \rightarrow action-classes
- ▶ `scores_to_action`
score predicted by the network \rightarrow action
- ▶ `cross_entropy_loss`
loss function: gt vs. prediction



d) Implement network

```
4  class ClassificationNetwork(torch.nn.Module):
5      def __init__(self):
6          """
7              1.1 d)
8              Implementation of the network layers. The image size of the input
9              observations is 96x96 pixels.
10             """
11             super().__init__()
12             gpu = torch.device('cuda')
```

- ▶ 2 to 3 convolution layers + ReLU
- ▶ 2 to 3 fully connected layers + ReLU
- ▶ Softmax

d) Implement network

```
torch.nn.Sequential(  
    torch.nn.Conv2d(in_channels, out_channels, filter_size, stride=*arg),  
    torch.nn.LeakyReLU(negative_slope=0.2))
```

```
torch.nn.Sequential(  
    torch.nn.Linear(in_size, out_size),  
    torch.nn.LeakyReLU(negative_slope=0.2))
```

- ▶ 2 to 3 convolution layers + ReLU
- ▶ 2 to 3 fully connected layers + ReLU
- ▶ Softmax

e) Forward pass, train and test

```
15     def forward(self, observation):
16         """
17         1.1 e)
18         The forward pass of the network. Returns the prediction for the given
19         input observation.
20         observation: torch.Tensor of size (batch_size, 96, 96, 3)
21         return       torch.Tensor of size (batch_size, C)
22         """
23     pass
```

- ▶ color channels or gray-scale
- ▶ `python3 main.py train`
- ▶ `python3 main.py test`
- ▶ hyper-parameter tuning

e) Forward pass, train and test

```
python3 main.py test
```


```
15     def evaluate():
16         """
17         """
18         infer_action = torch.load(trained_network_file)
19         infer_action.eval()
20         env = gym.make('CarRacing-v0')
21         gpu = torch.device('cuda')
22
23         for episode in range(5):
24             observation = env.reset()
25
26             reward_per_episode = 0
27             for t in range(500):
28                 env.render()
29                 action_scores = infer_action(
30                     torch.Tensor(np.ascontiguousarray(observation[None])).to(gpu))
31
32                 steer, gas, brake = infer_action.scores_to_action(action_scores)
33                 observation, reward, done, info = env.step([steer, gas, brake])
34                 reward_per_episode += reward
35
36             print('episode %d \t reward %f' % (episode, reward_per_episode))
37
```

f) Record own imitations


```
65  def record_imitations(imitations_folder):
66      """
67      Function to record own imitations by driving the car in the gym car-racing
68      environment.
69      imitations_folder: python string, the path to where the recorded imitations
70                        are to be saved
71
72      The controls are:
73      arrow keys:      control the car; steer left, steer right, gas, brake
74      ESC:              quit and close
75      SPACE:           restart on a new track
76      TAB:              save the current run
77      """
78      env = gym.make('CarRacing-v0').env
79      status = ControlStatus()
80      total_reward = 0.0
81
82      while not status.quit:
83          observations = []
84          actions = []
```

f) Record own imitations

```
python3 main.py teach
```



```
python3 main.py test
```



1.2

Network Improvements

a) Observations




a) Observations

```
49 def extract_sensor_values(self, observation, batch_size):
50     """
51     observation:    python list of batch_size many torch.Tensors of size
52                     (96, 96, 3)
53     batch_size:    int
54     return         torch.Tensors of size (batch_size, 1),
55                     torch.Tensors of size (batch_size, 4),
56                     torch.Tensors of size (batch_size, 1),
57                     torch.Tensors of size (batch_size, 1)
58     """
59     speed_crop = observation[:, 84:94, 12, 0].reshape(batch_size, -1)
60     speed = speed_crop.sum(dim=1, keepdim=True) / 255
61     abs_crop = observation[:, 84:94, 18:25:2, 2].reshape(batch_size, 10, 4)
62     abs_sensors = abs_crop.sum(dim=1) / 255
63     steer_crop = observation[:, 88, 38:58, 1].reshape(batch_size, -1)
64     steering = steer_crop.sum(dim=1, keepdim=True)
65     gyro_crop = observation[:, 88, 58:86, 0].reshape(batch_size, -1)
66     gyroscope = gyro_crop.sum(dim=1, keepdim=True)
67     return speed, abs_sensors.reshape(batch_size, 4), steering, gyroscope
```


b) MultiClass Prediction

- ▶ 4 binary classes:
steer left, steer right, accelerate, brake
- ▶ multi-class network
- ▶ actions → action-classes
- ▶ sigmoid activation function
- ▶ new loss function




c) Classification vs. regression

Regression network

- ▶ Which loss function?
- ▶ Advantages / drawbacks compared to the classification networks?
- ▶ Is it a reasonable approach given our training data?

d) Data augmentation

Create more training data by synthetically modifying the data.



e) Fine-tuning

- ▶ different network architectures
- ▶ learning rate adaptation
- ▶ dropout-, batch- or instance normalization
- ▶ different optimizers
- ▶ class imbalance

- ▶ try at least two ideas
- ▶ explain the motivation and the outcome

Competition

Competition

main.py

```
42 def calculate_score_for_leaderboard():
43     """
44     Evaluate the performance of the network. This is the function to be used for
45     the final ranking on the course-wide leader-board, only with a different set
46     of seeds. Better not change it.
47     """
48     infer_action = torch.load(trained_network_file, map_location='cpu')
49     infer_action.eval()
50     env = gym.make('CarRacing-v0')
51     # you can set it to torch.device('cuda') in case you have a gpu
52     device = torch.device('cpu')
53
54     seeds = [22597174, 68545857, 75568192, 91140053, 86018367,
55             49636746, 66759182, 91294619, 84274995, 31531469]
56     total_reward = 0
57
58     for episode in range(10):
59         env.seed(seeds[episode])
60         observation = env.reset()
61
62         reward_per_episode = 0
63         for t in range(600):
```

Questions?